Create a simple web application with multiple web servers and implement load balancing using AWS services. The project will involve setting up EC2 instances, configuring an Elastic Load Balancer (ELB),testing the load balancing functionality, autoscaling the servers

1.Launching Two EC2 Instances

Here we are going to launch two Ec2 Instances in US-east-1a & Us-east-1b

EC2 Specifications

Name-Webserver 1

Name-Webserver2

AMI-Amazon Linux 2023

Instance type-t2.micro

Key pair- p11.pem for Webserver1

Key pair- p12.pem for Webserver2

Networking

VPC- Default VPC

Subnet – Subnet in us-east-1a for Websever1

Subnet – Subnet in us-east-1b for Websever2

Security group

Create an SG which has SSH and HTTP access and attach it to both the instances.

Storage-

Keep default storage = 8GB as we are launching Linux instance.

Once launched wait till passing 2/2 check

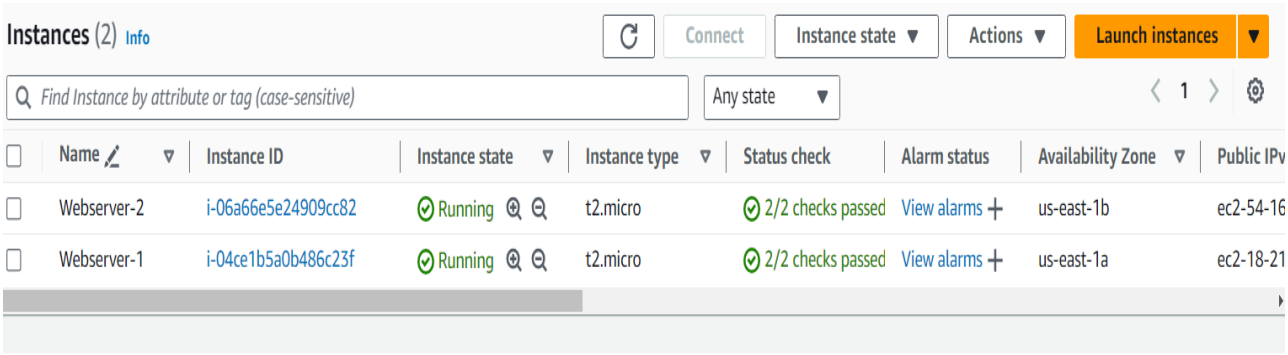

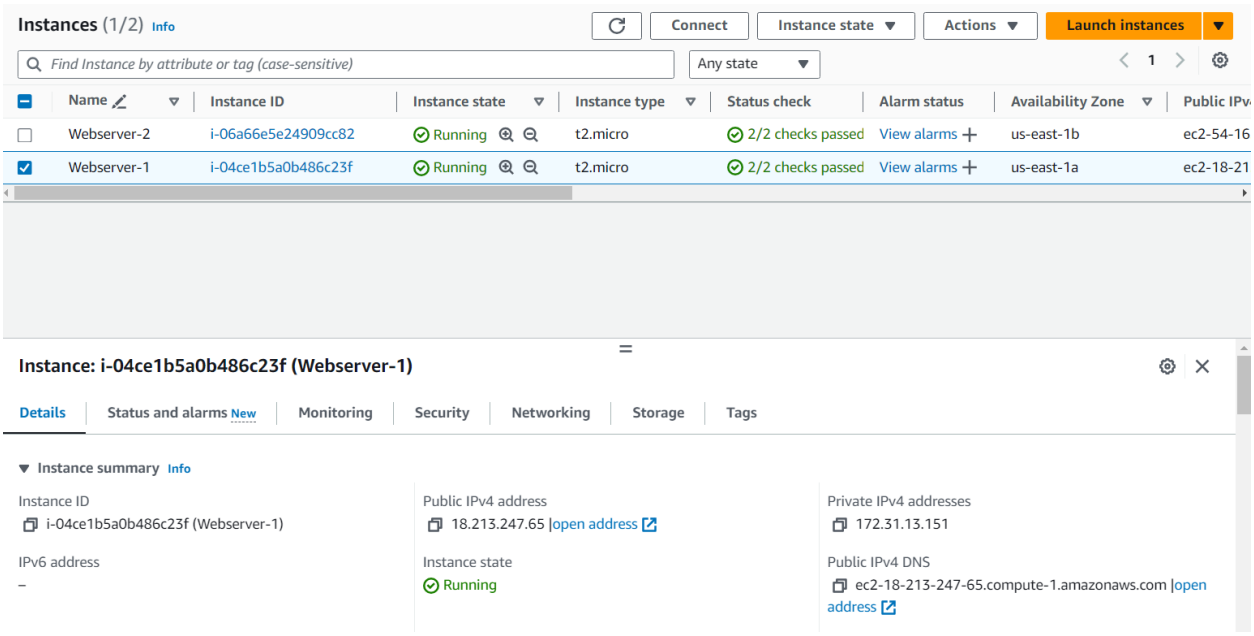

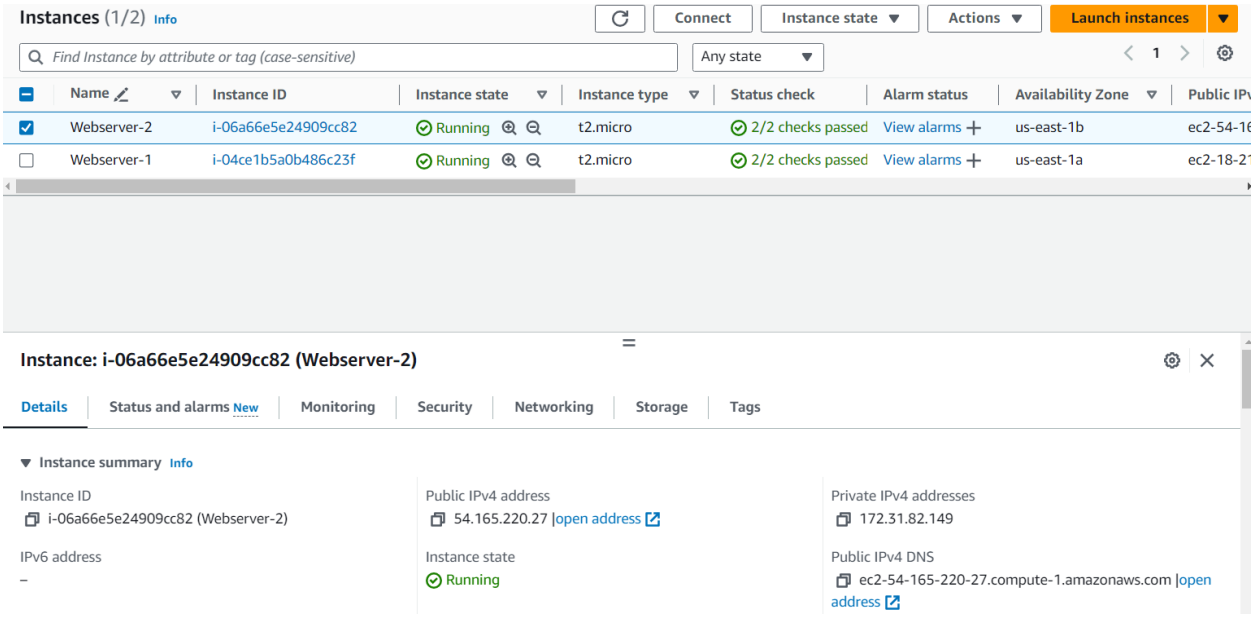

Security group

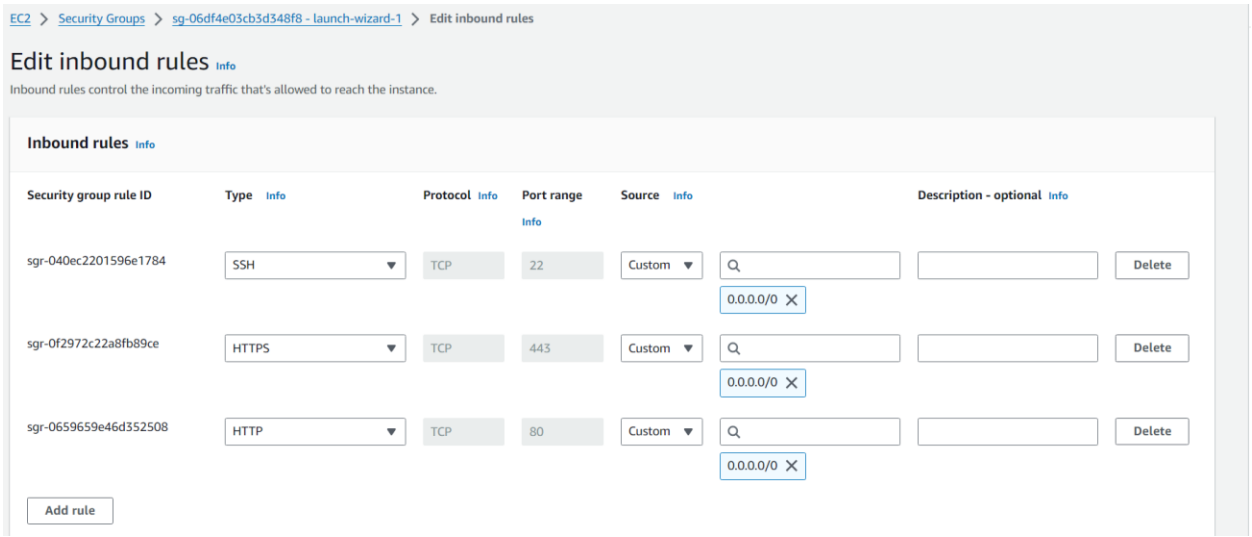

SSH to Webserver 1 & install Apache webserver on it

chmod 400 "p11.pem"

ssh -i "p11.pem" ec2-user@ec2-18-213-247-65.compute-1.amazonaws.com

you will get access to Webserver-1

```
hp@DESKTOP-UCM99CO MINGW64 /g/Cloud Project
$ ls
'AWS Project- CP.pdf'            Project1.pem    drive-download-20240113T024631Z-001.zip  'ec2-user@54.92.247.159'   my_web_app/   p12.pem   simple.zip      wordpress1.zip
'AWS Project-Architecting.pdf'   deploy.sh.txt  'ec2-user@35.175.139.17'/                 img/                       p11.pem       report/   winkel-wpcom/

hp@DESKTOP-UCM99CO MINGW64 /g/Cloud Project
$ chmod 400 "p11.pem"

hp@DESKTOP-UCM99CO MINGW64 /g/Cloud Project
$ ssh -i "p11.pem" ec2-user@ec2-18-213-247-65.compute-1.amazonaws.com
The authenticity of host 'ec2-18-213-247-65.compute-1.amazonaws.com (64:ff9b::12d5:f741)' can't be established.
ED25519 key fingerprint is SHA256:cZOraCJCSkG9pjCEqvq8T/h+oYuwwbufxGZCOsMpOOM.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-18-213-247-65.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
   ,     #_
   ~\_  ####_        Amazon Linux 2023
  ~~  \_#####\
  ~~     \###|
  ~~       \#/ ___   https://aws.amazon.com/linux/amazon-linux-2023
   ~~       V~' '->
    ~~~         /
      ~~._.   _/
         _/ _/
       _/m/'
[ec2-user@ip-172-31-13-151 ~]$
```

Execute Following commands on Webserver1

sudo su

yum update –y

```
       _/m/'
[ec2-user@ip-172-31-13-151 ~]$ sudo su
[root@ip-172-31-13-151 ec2-user]# yum update -y
Last metadata expiration check: 0:19:12 ago on Wed Jan 24 01:27:15 2024.
Dependencies resolved.
Nothing to do.
Complete!
[root@ip-172-31-13-151 ec2-user]#
```

yum install -y httpd

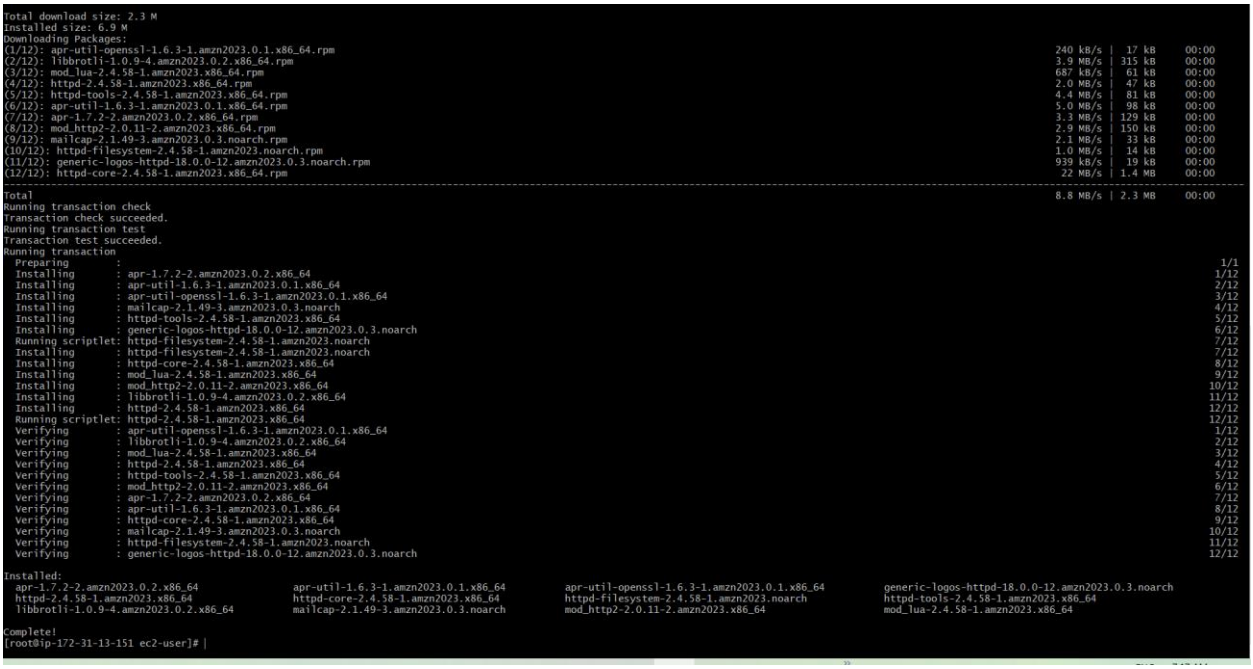

mkdir temp

cd temp

chown ec2-user:ec2-user temp/

chmod 700 temp/

exit

exit

scp -i p11.pem simple.zip ec2-user@18.213.247.65:/home/ec2-user/temp/

login again to webserver

ssh -i "p11.pem" ec2-user@ec2-18-213-247-65.compute-1.amazonaws.com

cd temp

you will find sample.zip file

unzip simple.zip

cd simple

mv * /var/www/html

systemctl enable httpd

systemctl start httpd

copy public ip of webserver1 and see whether webpage is accessible or not

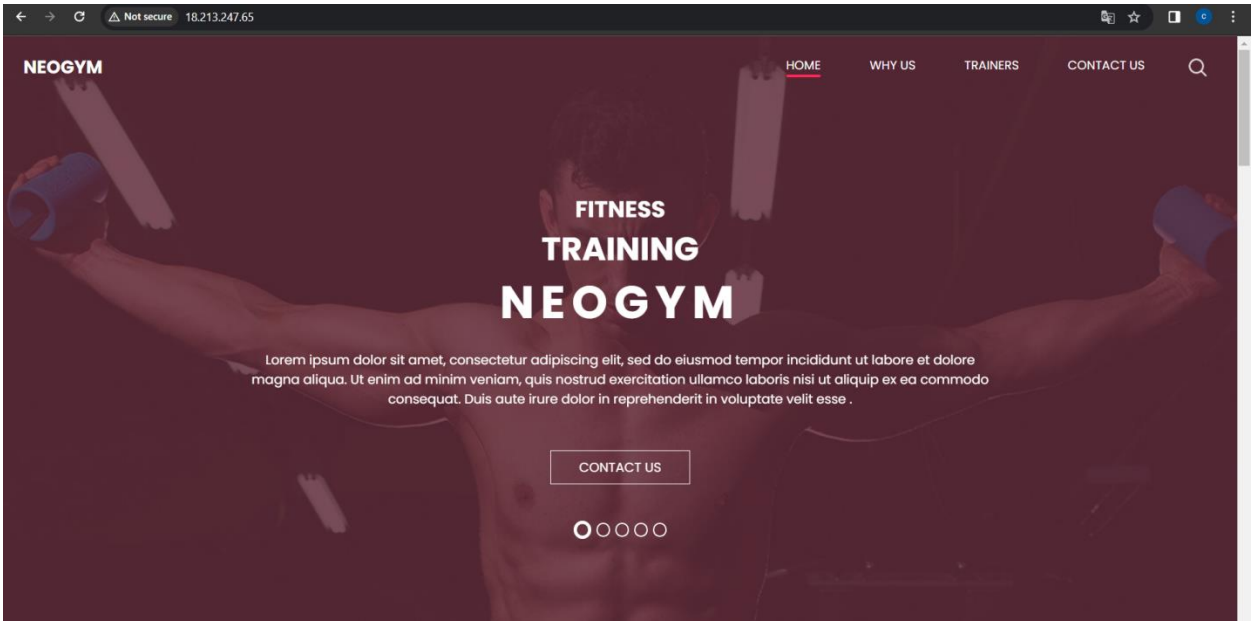

We successfully installed web application on webserver1

Repeat this process to install web application on Webserver2

chmod 400 "p12.pem"

ssh -i "p12.pem" ec2-user@ec2-54-165-220-27.compute-1.amazonaws.com

rest of the procedure will be same as per Webserver1

As we just started httpd without installing web app it looks like

It works!

scp -i p12.pem simple.zip ec2-user@54.165.220.27:/home/ec2-user/temp/

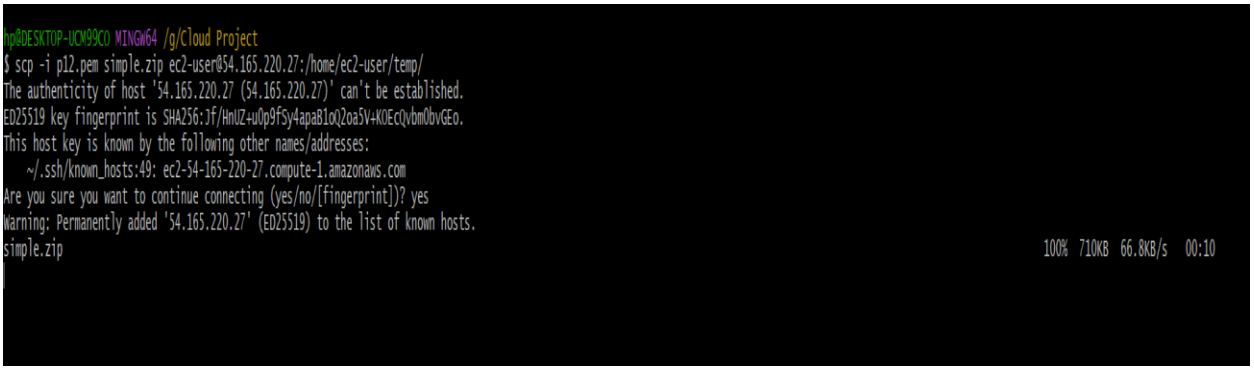

Log in to webserver2

sudo su

cd temp/

unzip simple.zip

cd simple

mv * /var/www/html

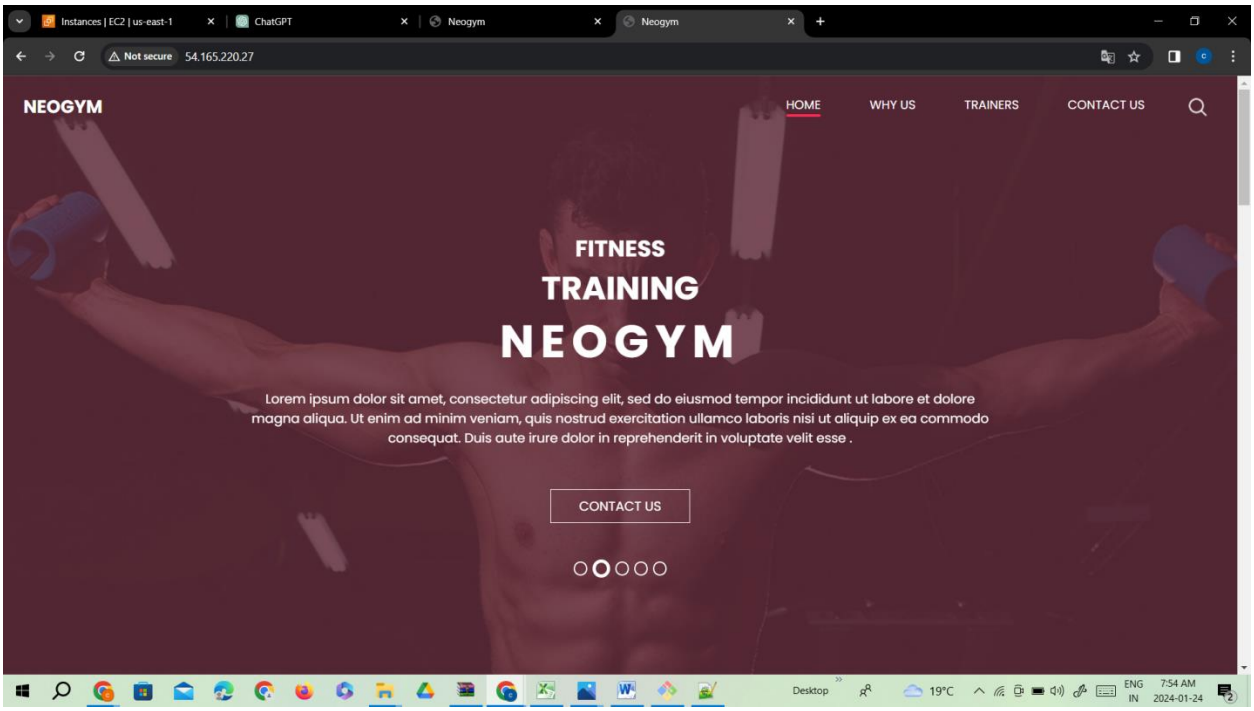

Configuring load balancer

Create load balancer

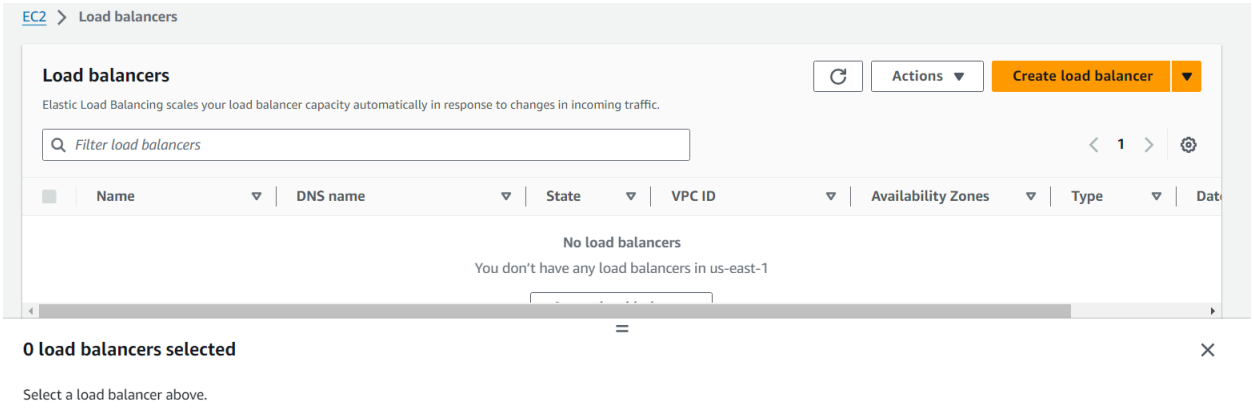

Create an Application load balancer

Load balancer name- ALB-001

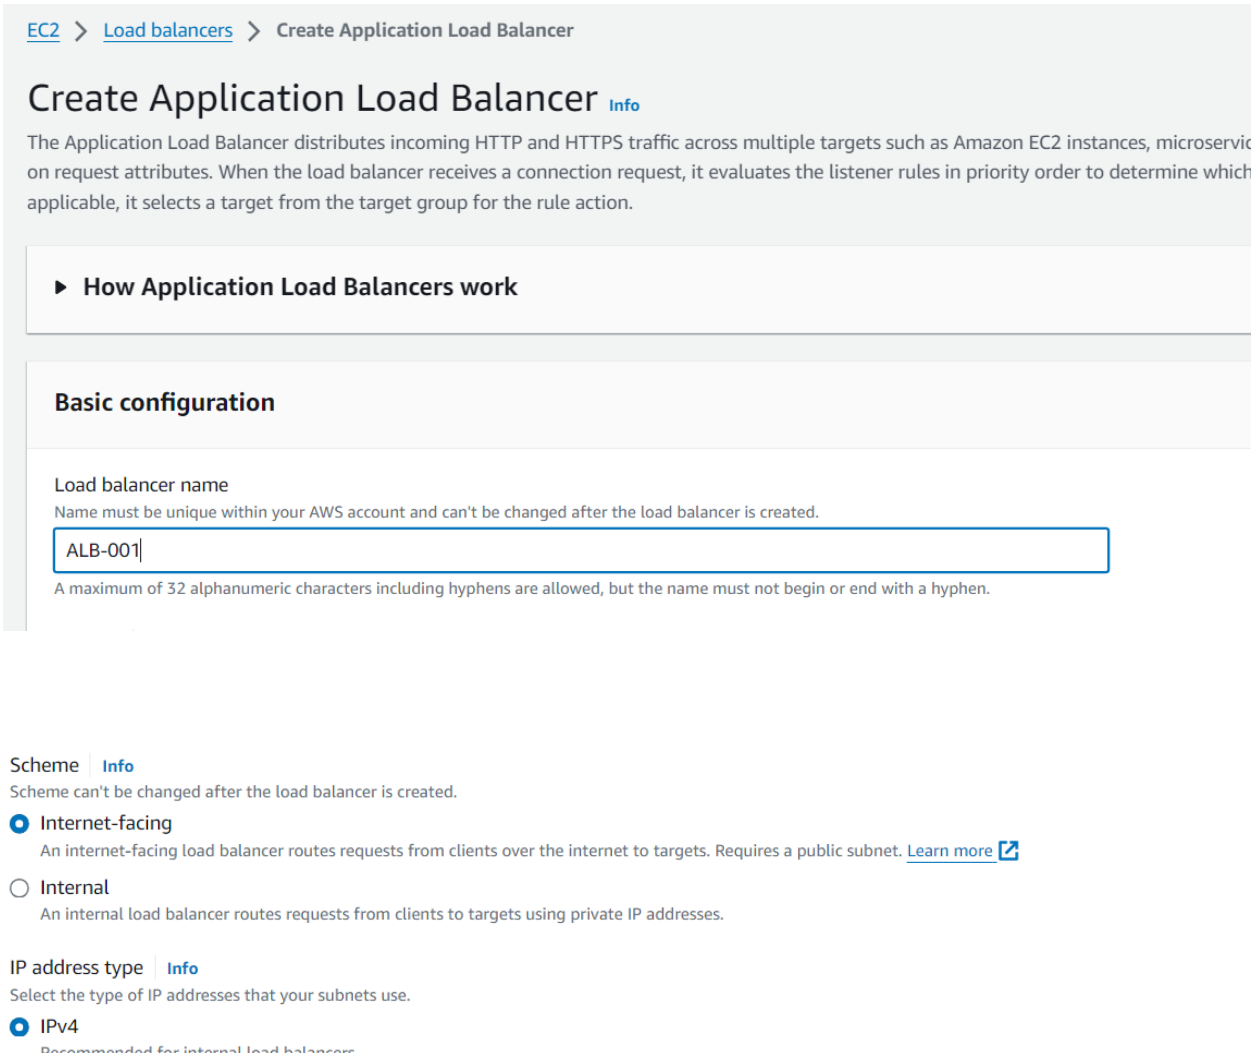
EC2 > Load balancers > Create Application Load Balancer
Create Application Load Balancer Info
The Application Load Balancer distributes incoming HTTP and HTTPS traffic across multiple targets such as Amazon EC2 instances, microservic
on request attributes. When the load balancer receives a connection request, it evaluates the listener rules in priority order to determine which
applicable, it selects a target from the target group for the rule action.
How Application Load Balancers work
Basic configuration
Load balancer name
Name must be unique within your AWS account and can't be changed after the load balancer is created.
ALB-001
A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Scheme Info
Scheme can't be changed after the load balancer is created.
Internet-facing
An internet-facing load balancer routes requests from clients over the internet to targets. Requires a public subnet. Learn more
Internal
An internal load balancer routes requests from clients to targets using private IP addresses.
IP address type Info
Select the type of IP addresses that your subnets use.
IPv4

Mappings Info

Select at least two Availability Zones and one subnet per zone. The load balancer routes traffic to targets in these Availability Zones only. Availability Zones th
balancer or the VPC are not available for selection.

us-east-1a (use1-az1)

Subnet

subnet-0c514ef92b26555f6

IPv4 address

Assigned by AWS

us-east-1b (use1-az2)

Subnet

subnet-03296ac2201e25af3

IPv4 address

Security groups Info

A security group is a set of firewall rules that control the traffic to your load balancer. Select an existing security group, or you can create a new security group.

Security groups

Select up to 5 security groups

ALB-SG
sg-06ecc0462e564f35b VPC: vpc-0b579b0189b65a94d

default
sg-09d2264ad27851504 VPC: vpc-0b579b0189b65a94d

Create Target group

Specify group details

Your load balancer routes requests to the targets in a target group and performs health checks on the targets.

Basic configuration

Settings in this section can't be changed after the target group is created.

Choose a target type

- (●) Instances
  - Supports load balancing to instances within a specific VPC.
  - Facilitates the use of Amazon EC2 Auto Scaling to manage and scale your EC2 capacity.
- ( ) IP addresses
  - Supports load balancing to VPC and on-premises resources.
  - Facilitates routing to multiple IP addresses and network interfaces on the same instance.
  - Offers flexibility with microservice based architectures, simplifying inter-application communication.
  - Supports IPv6 targets, enabling end-to-end IPv6 communication, and IPv4-to-IPv6 NAT.
- ( ) Lambda function
  - Facilitates routing to a single Lambda function.

Target group name

tg-001

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Protocol : Port

Choose a protocol for your target group that corresponds to the Load Balancer type that will route traffic to it. Some protocols now include anomaly detection for the targets and you can set mitigation options once your target group is created. This choice cannot be changed after creation

HTTP

80

1-65535

IP address type

Only targets with the indicated IP address type can be registered to this target group.

- (•) IPv4
  Each instance has a default network interface (eth0) that is assigned the primary private IPv4 address. The instance's primary private IPv4 address is the one that will be applied to the target.
- ( ) IPv6
  Each instance you register must have an assigned primary IPv6 address. This is configured on the instance's default network interface (eth0). Learn more

VPC

Select the VPC with the instances that you want to include in the target group. Only VPCs that support the IP address type selected above are available in this list.

-
vpc-0b579b0189b65a94d
IPv4: 172.31.0.0/16

Protocol version

- (•) HTTP1

Health check protocol

HTTP

Health check path

Use the default path of "/" to perform health checks on the root, or specify a custom path if preferred.

/

Up to 1024 characters allowed.

▼ **Advanced health check settings**

Restore defaults

Health check port

The port the load balancer uses when performing health checks on targets. By default, the health check port is the same as the target group's traffic port. However, you can specify a different port as an override.

- (●) Traffic port
- ( ) Override

Healthy threshold

The number of consecutive health checks successes required before considering an unhealthy target healthy.

5

2-10

Unhealthy threshold

The number of consecutive health check failures required before considering a target unhealthy.

2

Healthy threshold

The number of consecutive health checks successes required before considering an unhealthy target healthy.

5

2-10

Unhealthy threshold

The number of consecutive health check failures required before considering a target unhealthy.

2

2-10

Timeout

The amount of time, in seconds, during which no response means a failed health check.

5 seconds

2-120

Interval

The approximate amount of time between health checks of an individual target

10 seconds

5-300

Success codes

The HTTP codes to use when checking for a successful response from a target. You can specify multiple values (for example, "200,202") or a range of values (for example, "200-299").

200

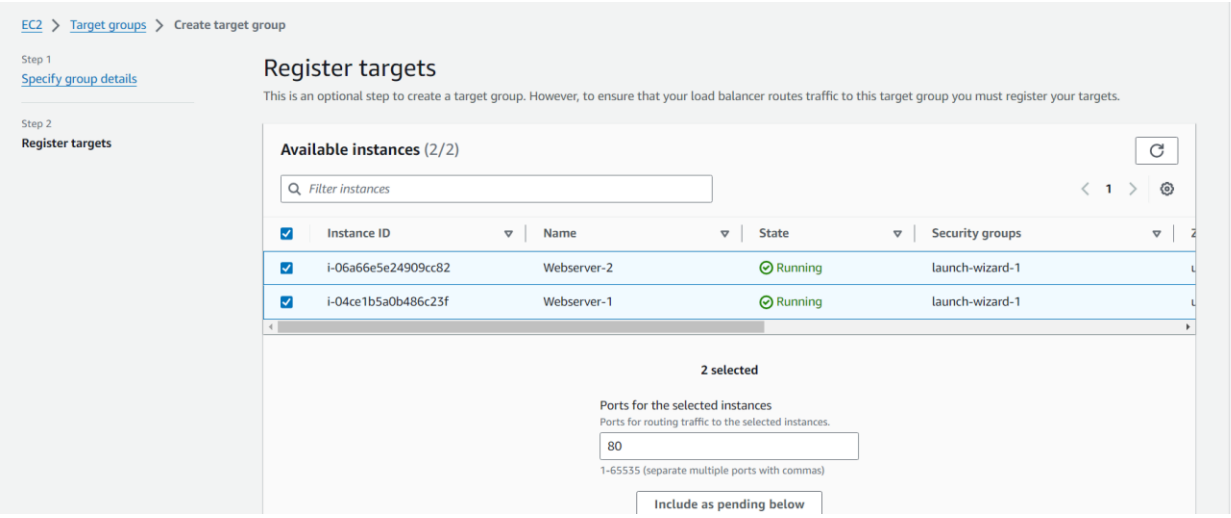

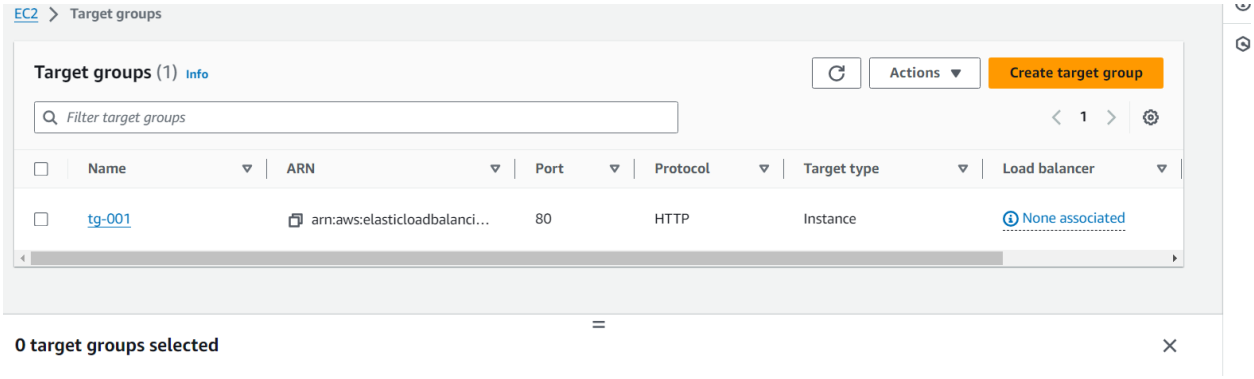

Listeners and routing Info

A listener is a process that checks for connection requests using the port and protocol you configure. The rules that you define for a listener determine how the load balancer routes reques to its registered targets.

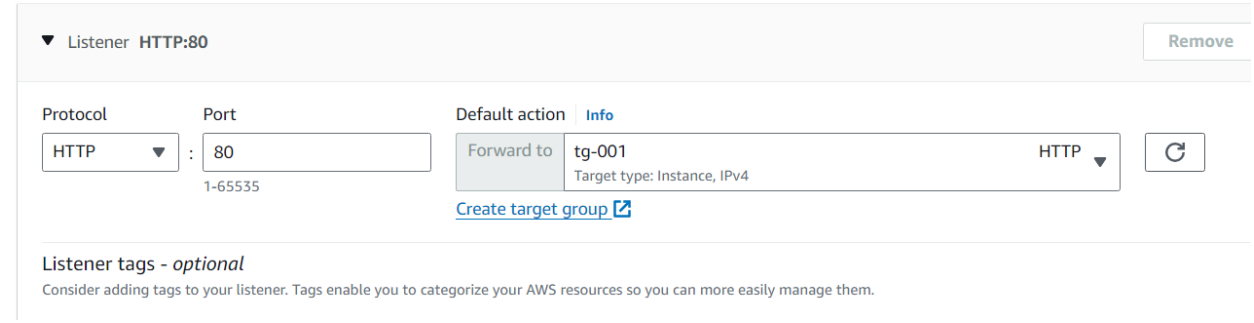

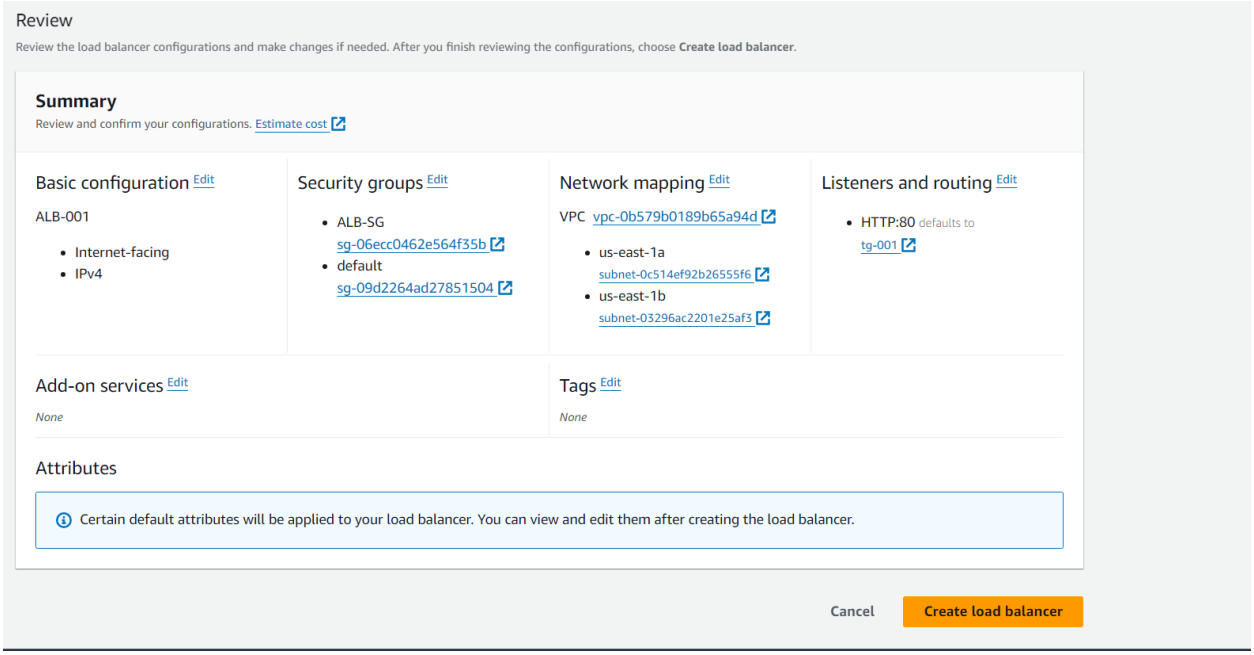

After creating load balancer wail till state is active

Once it is active

Copy ALB DNS name

Browse

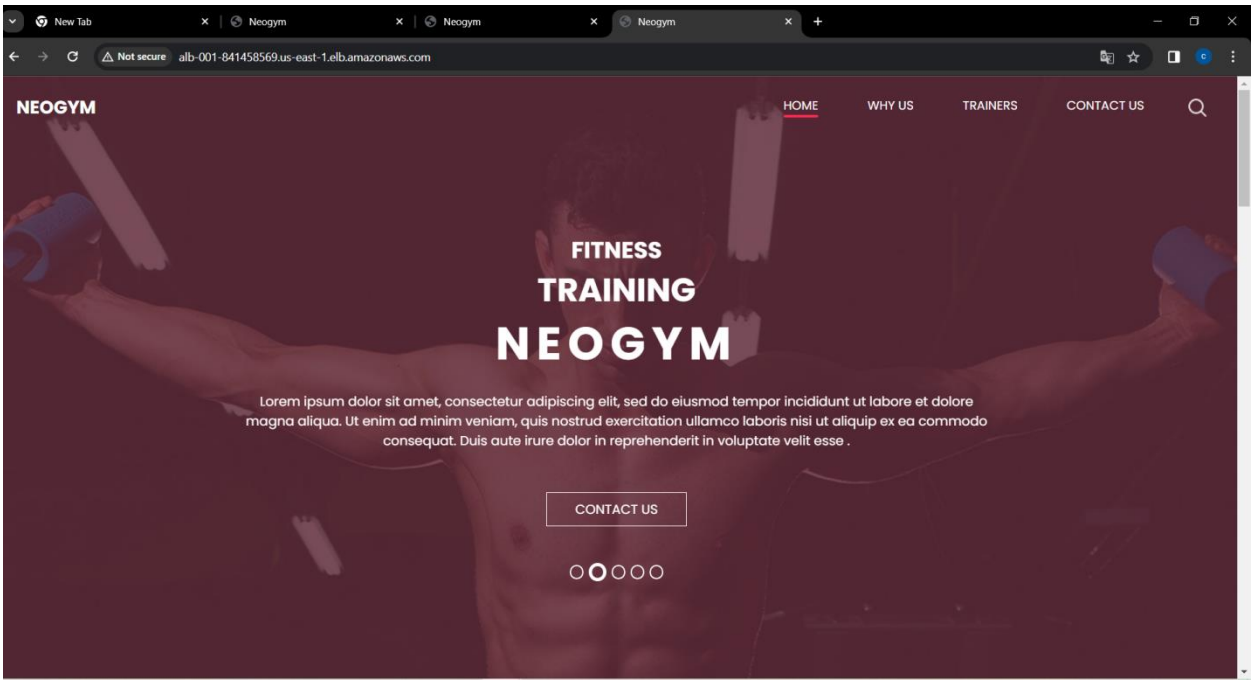

Auto Scaling group

To create Auto scaling group we have to create launch template

Here we stored application data in s3 bucket

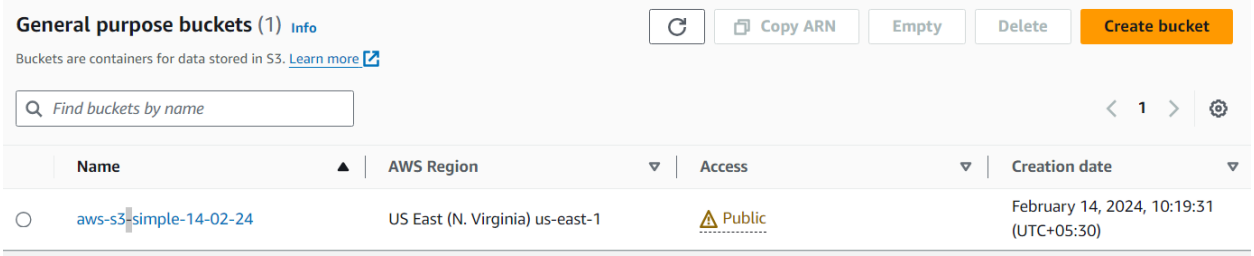

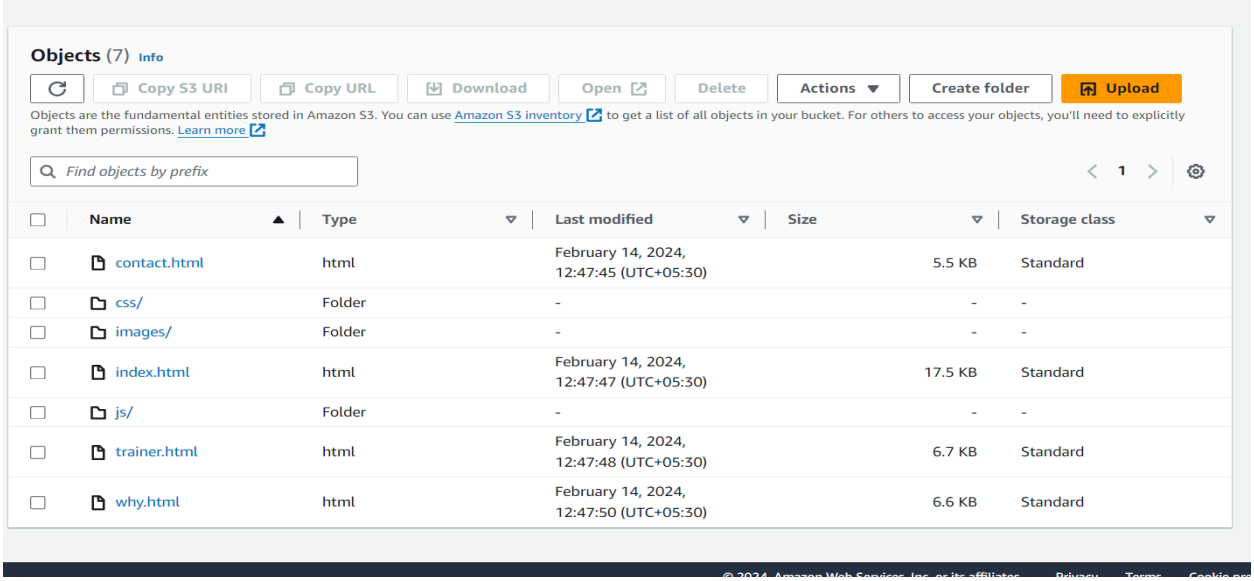

Launch Template

EC2 > Launch templates > Create launch template

Create launch template

Creating a launch template allows you to create a saved instance configuration that can be reused, shared and launched a later time. Templates can have multiple versions.

Launch template name and description

Launch template name - *required*

Sample-Web-APP

Must be unique to this account. Max 128 chars. No spaces or special characters like '&', '*', '@'.

Template version description

Template is to deploy application in production enviornment

Max 255 chars

Application and OS Images (Amazon Machine Image) - required Info

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Search our full catalog including 1000s of application and OS images

Recents | Quick Start

Amazon Linux | macOS | Ubuntu | Windows | Red Hat | SUSE Li

Browse more AMIs

Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

Amazon Linux 2023 AMI — Free tier eligible

ami-0e731c8a588258d0d (64-bit (x86), uefi-preferred) / ami-0bbebc09f0a12d4d9 (64-bit (Arm), uefi)

Virtualization: hvm ENA enabled: true Root device type: ebs

Description

Amazon Linux 2023 AMI 2023.3.20240205.2 x86_64 HVM kernel-6.1

Instance type Info | Get advice

Advanced

Instance type

t2.micro — Free tier eligible

Family: t2 1 vCPU 1 GiB Memory Current generation: true

On-Demand Windows base pricing: 0.0162 USD per Hour

On-Demand SUSE base pricing: 0.0116 USD per Hour

On-Demand RHEL base pricing: 0.0716 USD per Hour

On-Demand Linux base pricing: 0.0116 USD per Hour

All generations

Compare instance types

Additional costs apply for AMIs with pre-installed software

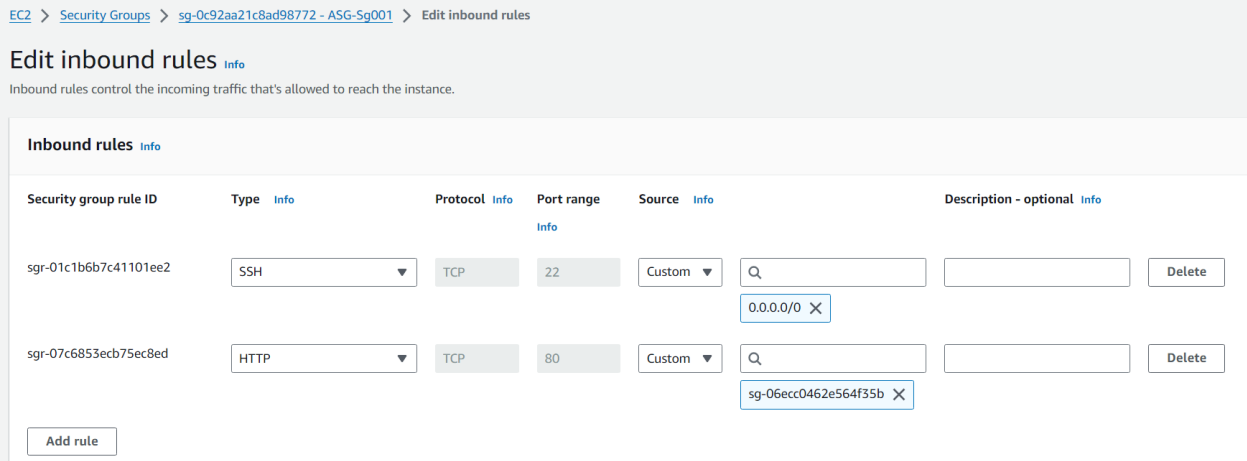

- **sg-06ecc0462e564f35b** is a security group for load balancer which allows traffic from internet.

  Here we allow traffic from load balancer to our EC2 instances which are created from auto scaling group.

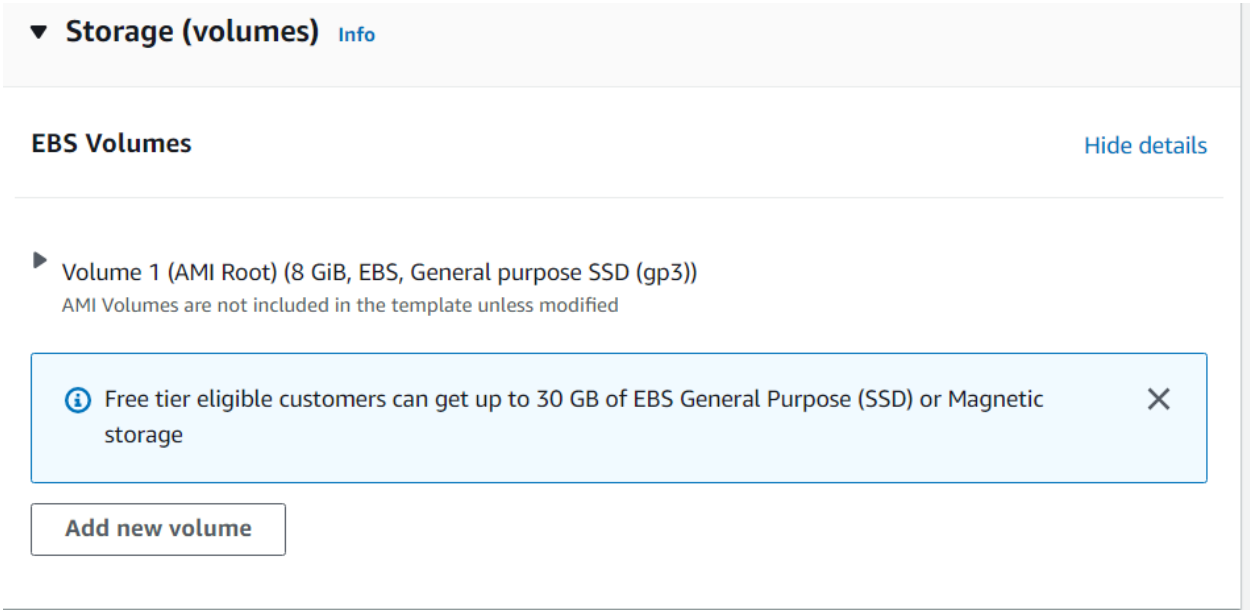

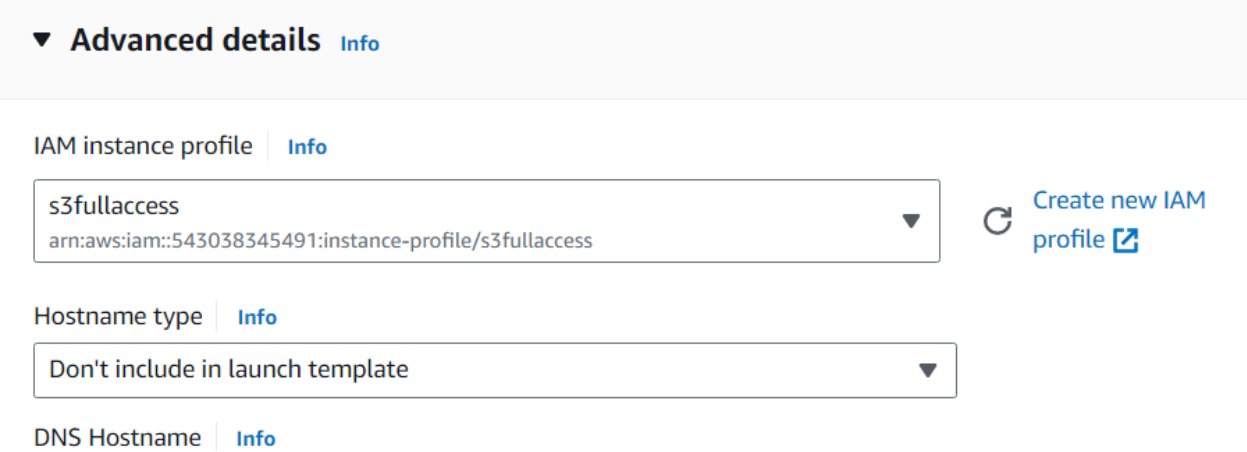

This roles gives permission to ec2 instance to communicate with s3 & access data from s3.

```
{
   "Version": "2012-10-17",
   "Statement": [
      {
         "Effect": "Allow",
         "Action": [
            "s3:*",
            "s3-object-lambda:*"
         ],
         "Resource": "*"
      }
   ]
}
```

We Have to add user data bash script

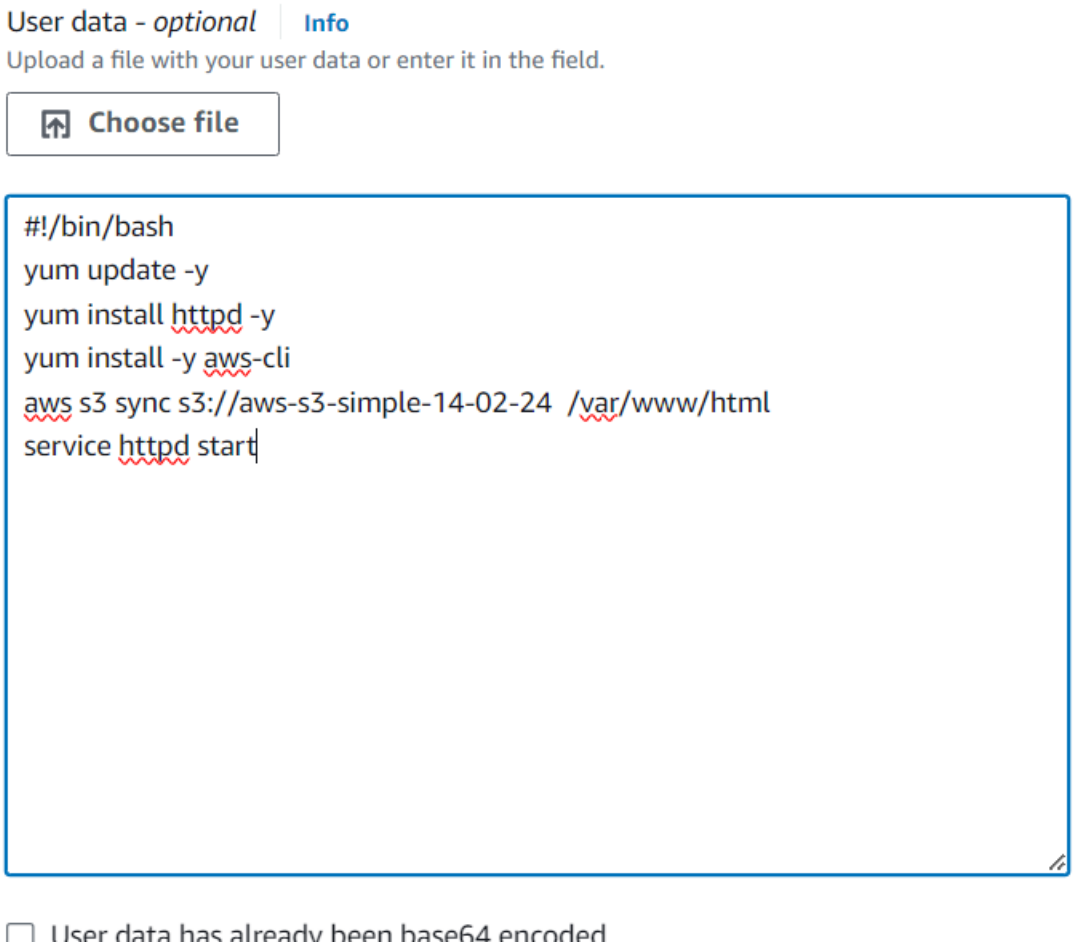


```
#!/bin/bash
yum update -y
yum install httpd -y
yum install -y aws-cli
aws s3 sync s3://aws-s3-simple-14-02-24  /var/www/html
service httpd start
```

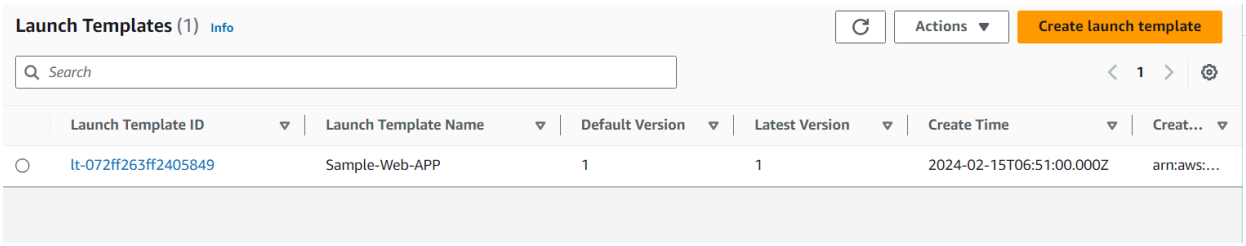

Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group.

Name

Auto Scaling group name

Enter a name to identify the group.

ASG-Sample-Web-app

Must be unique to this account in the current Region and no more than 255 characters.

Launch template Info

For accounts created after May 31, 2023, the EC2 console only supports creating Auto Scaling groups with launch templates. Creating Auto Scaling groups with launch configurations is not recommended but still available via the CLI and API until December 31, 2023.

Launch template

Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

Sample-Web-APP

Create a launch template

Version

Default (1)

Create a launch template version

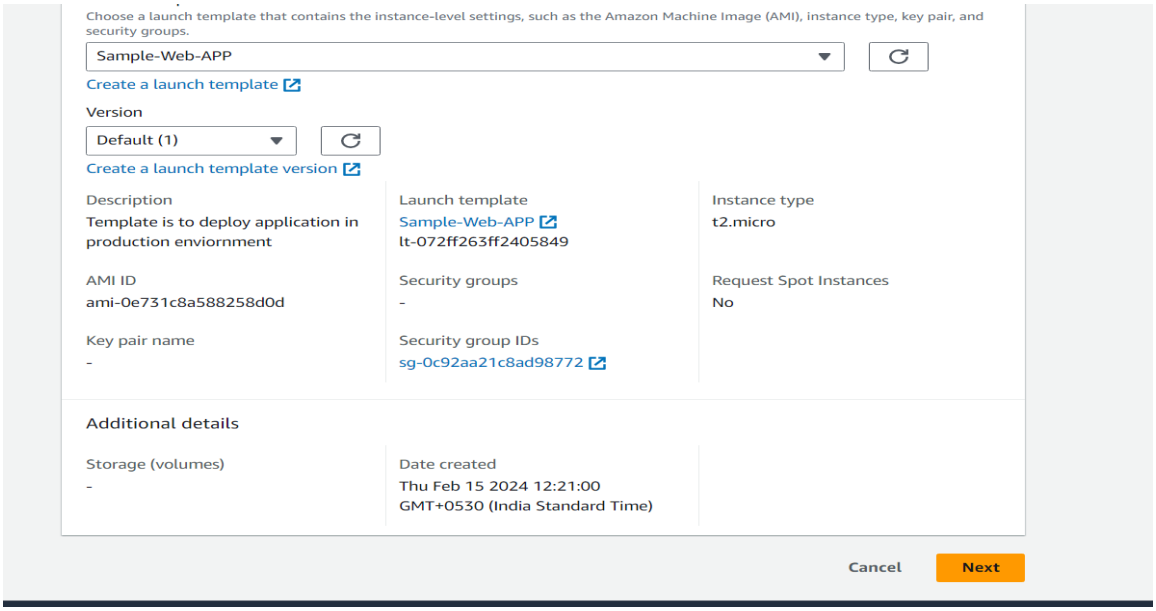
Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.
Sample-Web-APP
Create a launch template
Version
Default (1)
Create a launch template version
Description
Template is to deploy application in production enviornment
Launch template
Sample-Web-APP
lt-072ff263ff2405849
Instance type
t2.micro
AMI ID
ami-0e731c8a588258d0d
Security groups
-
Request Spot Instances
No
Key pair name
-
Security group IDs
sg-0c92aa21c8ad98772
Additional details
Storage (volumes)
-
Date created
Thu Feb 15 2024 12:21:00
GMT+0530 (India Standard Time)
Cancel
Next

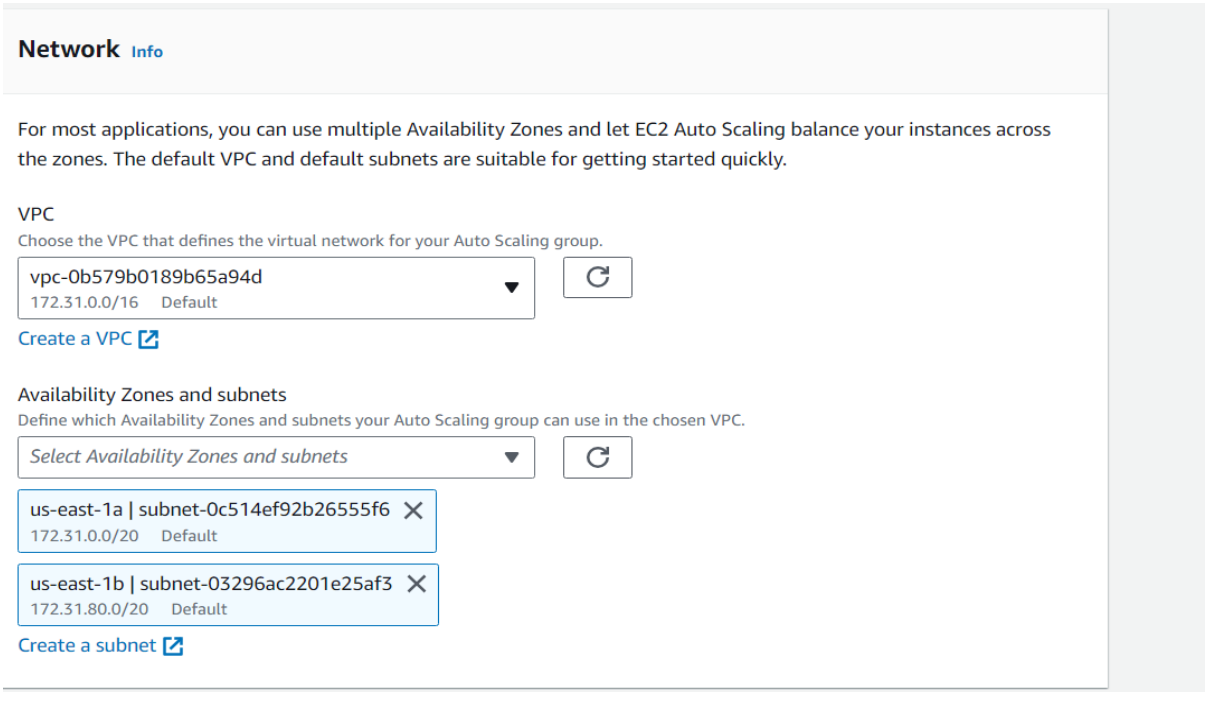
Network Info
For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.
VPC
Choose the VPC that defines the virtual network for your Auto Scaling group.
vpc-0b579b0189b65a94d
172.31.0.0/16 Default
Create a VPC
Availability Zones and subnets
Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.
Select Availability Zones and subnets
us-east-1a | subnet-0c514ef92b26555f6
172.31.0.0/20 Default
us-east-1b | subnet-03296ac2201e25af3
172.31.80.0/20 Default
Create a subnet

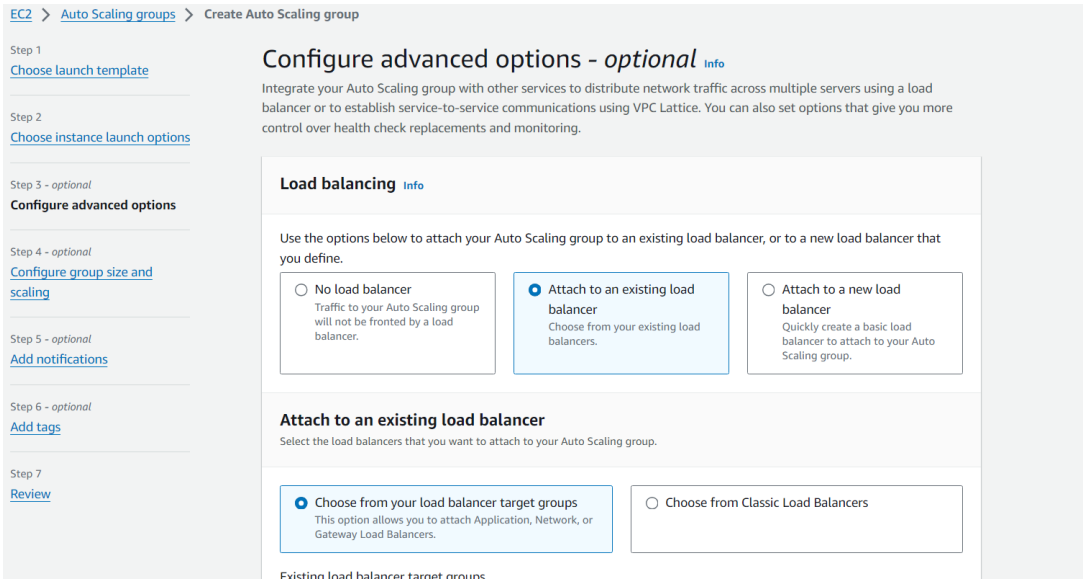
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1
Choose launch template
Step 2
Choose instance launch options
Step 3 - optional
Configure advanced options
Step 4 - optional
Configure group size and scaling
Step 5 - optional
Add notifications
Step 6 - optional
Add tags
Step 7
Review
Configure advanced options - optional Info
Integrate your Auto Scaling group with other services to distribute network traffic across multiple servers using a load balancer or to establish service-to-service communications using VPC Lattice. You can also set options that give you more control over health check replacements and monitoring.
Load balancing Info
Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.
No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.
Attach to an existing load balancer
Choose from your existing load balancers.
Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.
Attach to an existing load balancer
Select the load balancers that you want to attach to your Auto Scaling group.
Choose from your load balancer target groups
This option allows you to attach Application, Network, or Gateway Load Balancers.
Choose from Classic Load Balancers
Existing load balancer target groups

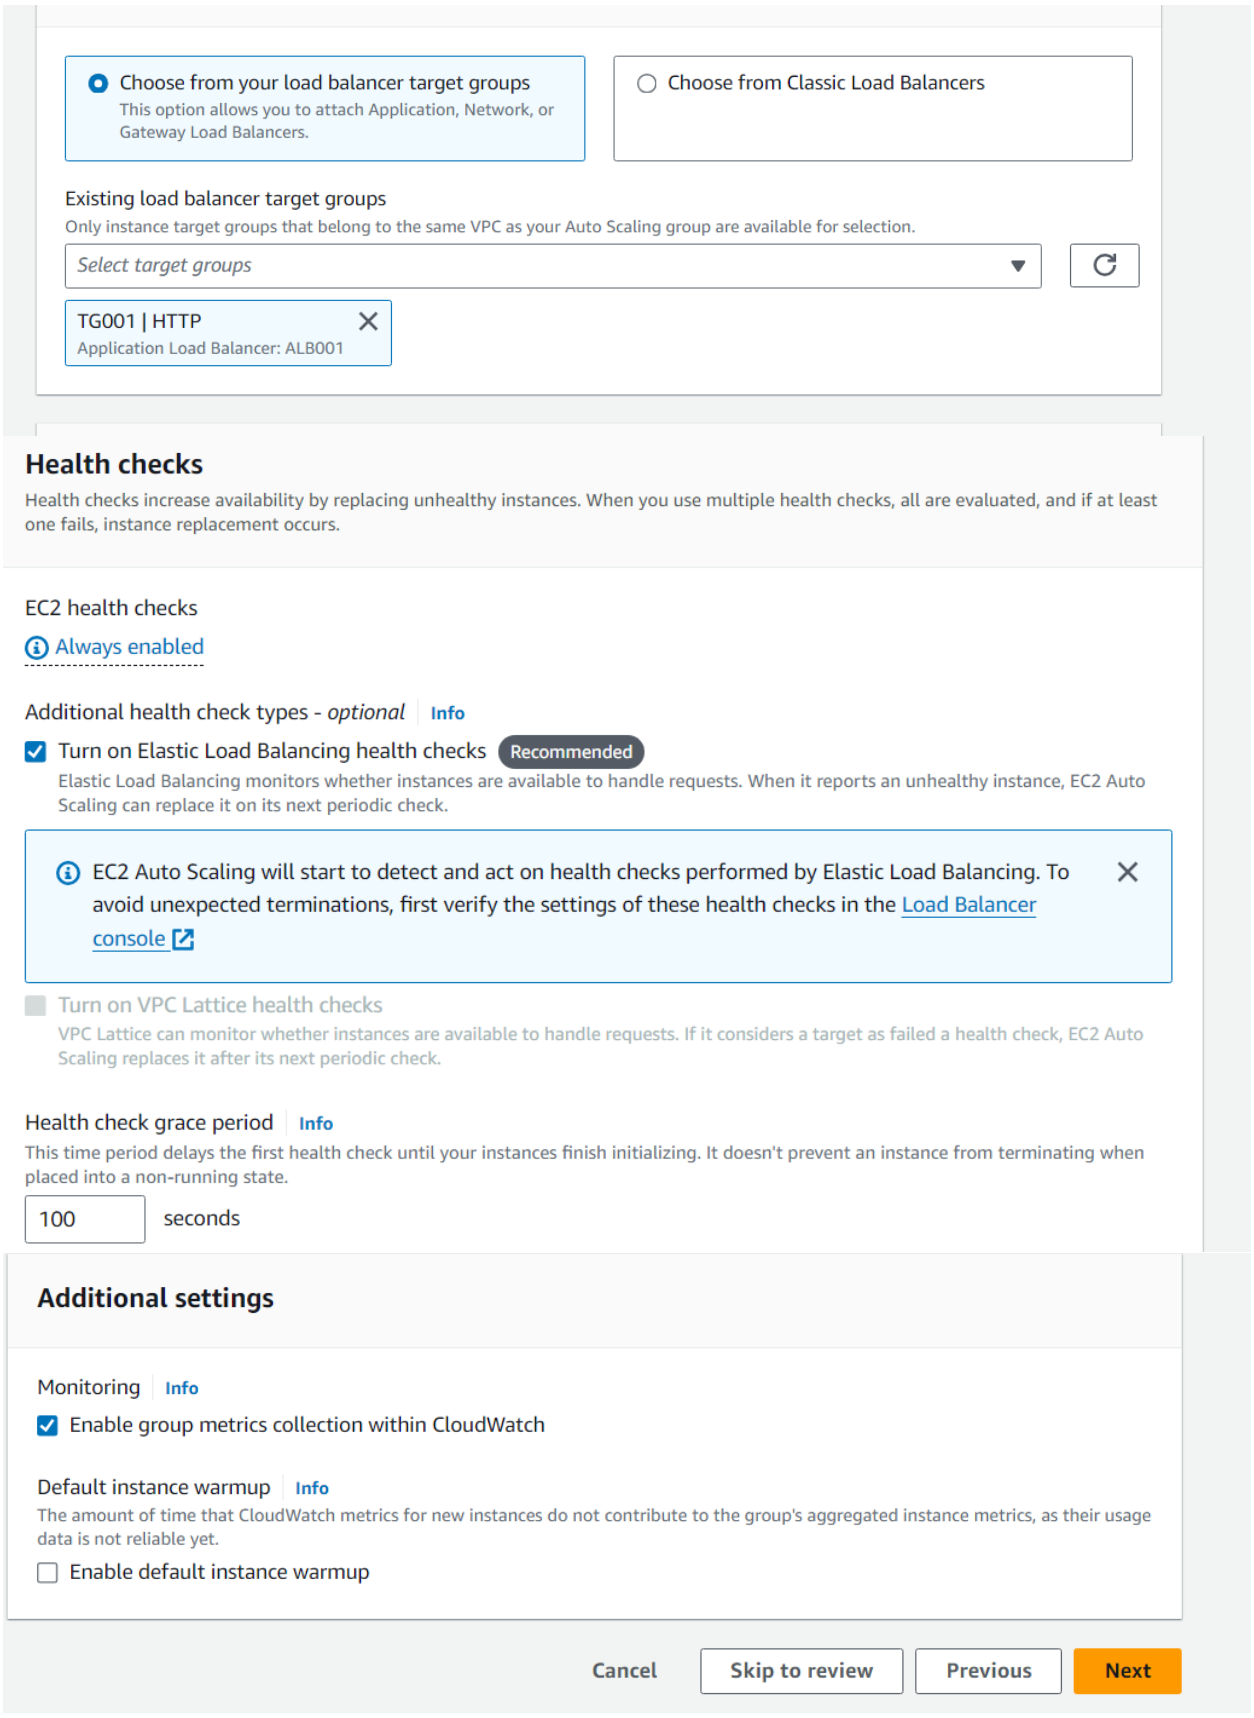
Choose from your load balancer target groups
This option allows you to attach Application, Network, or Gateway Load Balancers.
Choose from Classic Load Balancers
Existing load balancer target groups
Only instance target groups that belong to the same VPC as your Auto Scaling group are available for selection.
Select target groups
TG001 | HTTP
Application Load Balancer: ALB001
Health checks
Health checks increase availability by replacing unhealthy instances. When you use multiple health checks, all are evaluated, and if at least one fails, instance replacement occurs.
EC2 health checks
Always enabled
Additional health check types - optional
Info
Turn on Elastic Load Balancing health checks
Recommended
Elastic Load Balancing monitors whether instances are available to handle requests. When it reports an unhealthy instance, EC2 Auto Scaling can replace it on its next periodic check.
EC2 Auto Scaling will start to detect and act on health checks performed by Elastic Load Balancing. To avoid unexpected terminations, first verify the settings of these health checks in the Load Balancer console
Turn on VPC Lattice health checks
VPC Lattice can monitor whether instances are available to handle requests. If it considers a target as failed a health check, EC2 Auto Scaling replaces it after its next periodic check.
Health check grace period
Info
This time period delays the first health check until your instances finish initializing. It doesn't prevent an instance from terminating when placed into a non-running state.
100
seconds
Additional settings
Monitoring
Info
Enable group metrics collection within CloudWatch
Default instance warmup
Info
The amount of time that CloudWatch metrics for new instances do not contribute to the group's aggregated instance metrics, as their usage data is not reliable yet.
Enable default instance warmup
Cancel
Skip to review
Previous
Next

Configure group size and scaling - *optional* Info

Define your group's desired capacity and scaling limits. You can optionally add automatic scaling to adjust the size of your group.

Group size Info

Set the initial size of the Auto Scaling group. After creating the group, you can change its size to meet demand, either manually or by using automatic scaling.

Desired capacity type

Choose the unit of measurement for the desired capacity value. vCPUs and Memory(GiB) are only supported for mixed instances groups configured with a set of instance attributes.

Units (number of instances)

Desired capacity

Specify your group size.

2

Scaling Info

You can resize your Auto Scaling group manually or automatically to meet changes in demand.

Scaling limits

Set limits on how much your desired capacity can be increased or decreased.

Min desired capacity

2

Equal or less than desired capacity

Max desired capacity

4

Equal or greater than desired capacity

Automatic scaling - *optional*

Choose whether to use a target tracking policy Info

You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.

- ○ No scaling policies
  Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.
- ◉ Target tracking scaling policy
  Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.

Automatic scaling - *optional*

Choose whether to use a target tracking policy | Info

You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.

- ○ No scaling policies
  Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.
- ◉ Target tracking scaling policy
  Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.

Scaling policy name

Monitor CPU utilization metric

Metric type | Info

Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.

Average CPU utilization

Target value

80

Instance warmup | Info

100 seconds

☐ Disable scale in to create only a scale-out policy

Add notifications - *optional* Info

Send notifications to SNS topics whenever Amazon EC2 Auto Scaling launches or terminates the EC2 instances in your Auto Scaling group.

▼ Notification 1

Remove

Send a notification to

my-sns-topic-1

With these recipients

cpandharpure@gmail.com

Use existing topic

Event types

Notify subscribers whenever instances

- ☑ Launch
- ☑ Terminate
- ☑ Fail to launch
- ☑ Fail to terminate

Details

Type Info

Topic type cannot be modified after topic is created

- ○ FIFO (first-in, first-out)
  - Strictly-preserved message ordering
  - Exactly-once message delivery
  - High throughput, up to 300 publishes/second
  - Subscription protocols: SQS
- ◉ Standard
  - Best-effort message ordering
  - At-least once message delivery
  - Highest throughput in publishes/second
  - Subscription protocols: SQS, Lambda, HTTP, SMS, email, mobile application endpoints

Name

My-SNS-001

Maximum 256 characters. Can include alphanumeric characters, hyphens (-) and underscores (_).

Display name - *optional* Info

To use this topic with SMS subscriptions, enter a display name. Only the first 10 characters are displayed in an SMS message.

Notifications from the Autoscaling group

Maximum 100 characters.

Create subscription

Details

Topic ARN

arn:aws:sns:us-east-1:543038345491:My-SNS-001

Protocol

The type of endpoint to subscribe

Email

Endpoint

An email address that can receive notifications from Amazon SNS.

cpandharpure@gmail.com

After your subscription is created, you must confirm it. Info

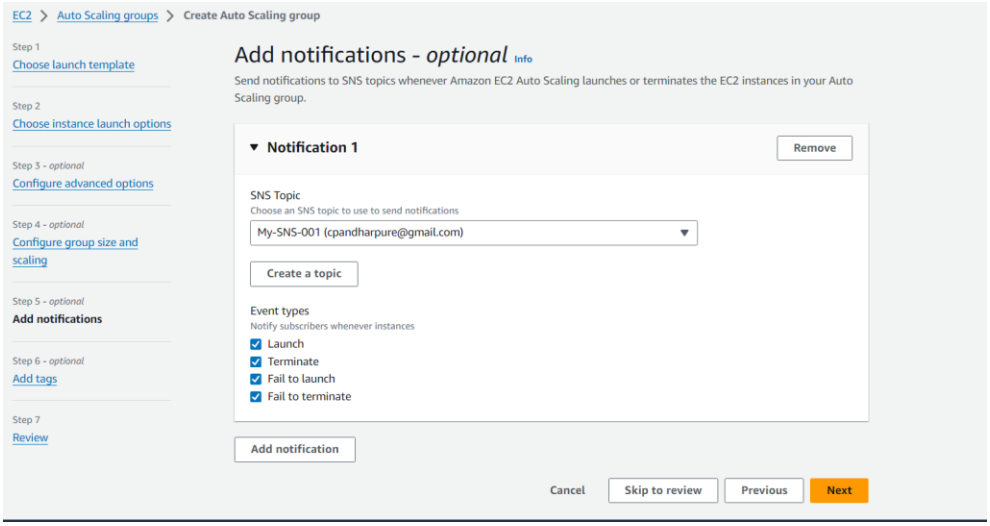
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1
Choose launch template
Step 2
Choose instance launch options
Step 3 - optional
Configure advanced options
Step 4 - optional
Configure group size and scaling
Step 5 - optional
Add notifications
Step 6 - optional
Add tags
Step 7
Review
Add notifications - optional Info
Send notifications to SNS topics whenever Amazon EC2 Auto Scaling launches or terminates the EC2 instances in your Auto Scaling group.
Notification 1
Remove
SNS Topic
Choose an SNS topic to use to send notifications
My-SNS-001 (cpandharpure@gmail.com)
Create a topic
Event types
Notify subscribers whenever instances
Launch
Terminate
Fail to launch
Fail to terminate
Add notification
Cancel
Skip to review
Previous
Next

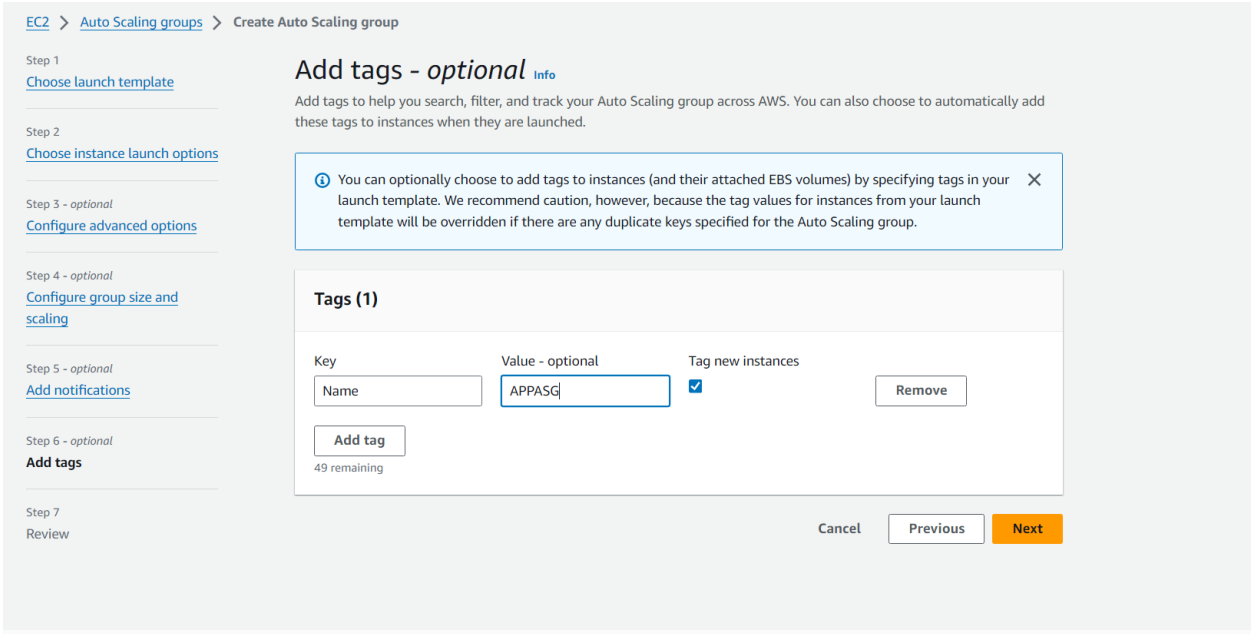
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1
Choose launch template
Step 2
Choose instance launch options
Step 3 - optional
Configure advanced options
Step 4 - optional
Configure group size and scaling
Step 5 - optional
Add notifications
Step 6 - optional
Add tags
Step 7
Review
Add tags - optional Info
Add tags to help you search, filter, and track your Auto Scaling group across AWS. You can also choose to automatically add these tags to instances when they are launched.
You can optionally choose to add tags to instances (and their attached EBS volumes) by specifying tags in your launch template. We recommend caution, however, because the tag values for instances from your launch template will be overridden if there are any duplicate keys specified for the Auto Scaling group.
Tags (1)
Key
Value - optional
Tag new instances
Name
APPASG
Remove
Add tag
49 remaining
Cancel
Previous
Next

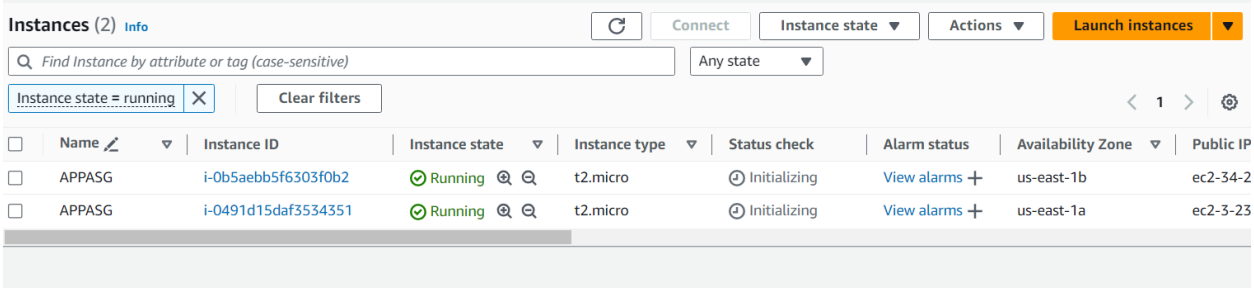
Instances (2) Info
Connect
Instance state
Actions
Launch instances
Find Instance by attribute or tag (case-sensitive)
Any state
Instance state = running
Clear filters
1
Name
Instance ID
Instance state
Instance type
Status check
Alarm status
Availability Zone
Public IP
APPASG
i-0b5aebb5f6303f0b2
Running
t2.micro
Initializing
View alarms
us-east-1b
ec2-34-2
APPASG
i-0491d15daf3534351
Running
t2.micro
Initializing
View alarms
us-east-1a
ec2-3-23

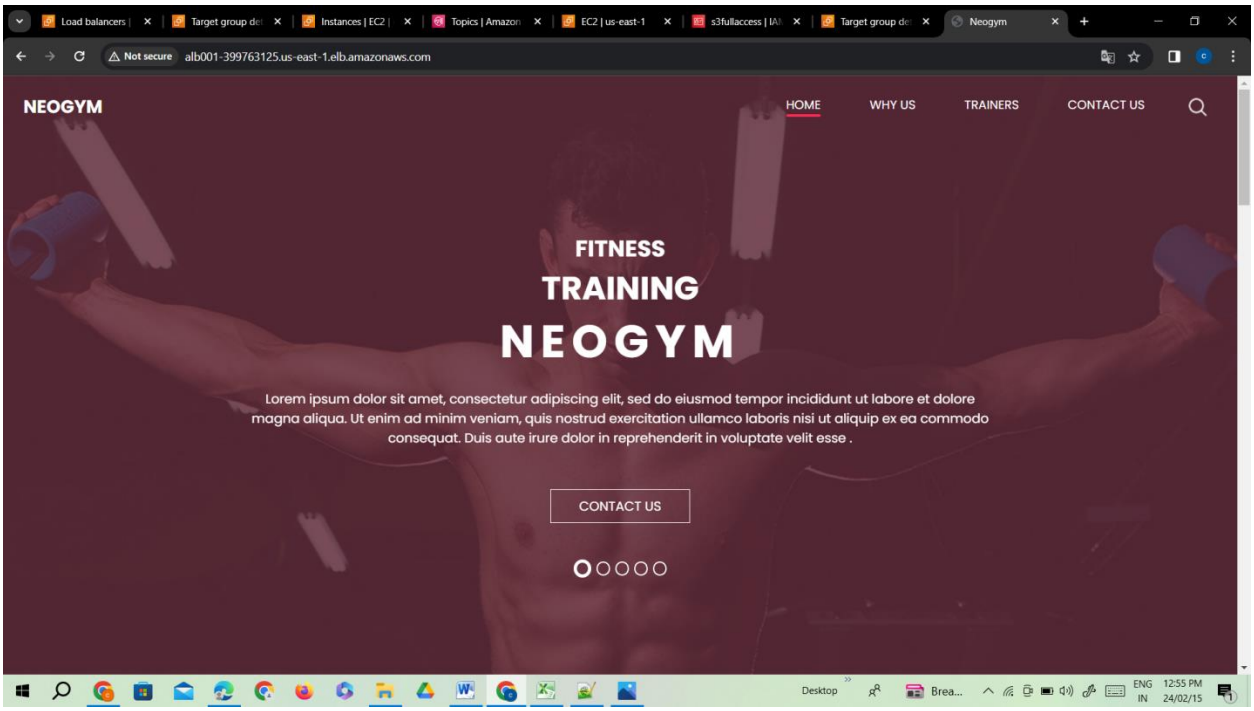
Not secure
alb001-399763125.us-east-1.elb.amazonaws.com
NEOGYM
HOME
WHY US
TRAINERS
CONTACT US
FITNESS
TRAINING
NEOGYM
Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse .
CONTACT US

0000

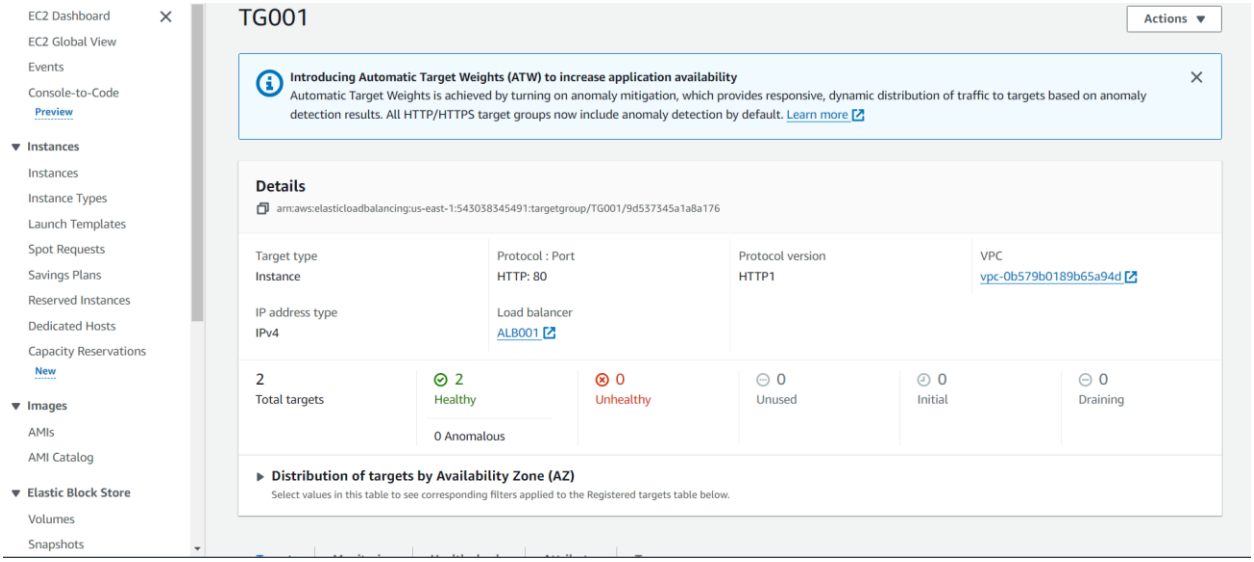
EC2 Dashboard
EC2 Global View
Events
Console-to-Code
Preview
Instances
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
New
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
TG001
Actions
Introducing Automatic Target Weights (ATW) to increase application availability
Automatic Target Weights is achieved by turning on anomaly mitigation, which provides responsive, dynamic distribution of traffic to targets based on anomaly detection results. All HTTP/HTTPS target groups now include anomaly detection by default. Learn more
Details
arn:aws:elasticloadbalancing:us-east-1:543038345491:targetgroup/TG001/9d537345a1a8a176
Target type
Instance
Protocol : Port
HTTP: 80
Protocol version
HTTP1
VPC
vpc-0b579b0189b65a94d
IP address type
IPv4
Load balancer
ALB001
2
Total targets
2
Healthy
0 Anomalous
0
Unhealthy
0
Unused
0
Initial
0
Draining
Distribution of targets by Availability Zone (AZ)
Select values in this table to see corresponding filters applied to the Registered targets table below.

Auto Scaling: launch for group "ASG-Sample-APP" Inbox ×

Notifications from the Autoscaling group 12:53 PM (6 minutes ago)
Service: AWS Auto Scaling Time: 2024-02-15T07:23:45.513Z RequestId: 140636ae-97a8-3bb9-e892-068db67512d3 Event: autoscaling:EC2_INSTANCE_LAUNCH AccountId: 54303

Notifications from the Autoscaling group <no-reply@sns.amazonaws.com> 12:53 PM (6 minutes ago)
to me

Service: AWS Auto Scaling
Time: 2024-02-15T07:23:45.649Z
RequestId: f6a636ae-97a9-71c3-7fd3-ccc662794ab6
Event: autoscaling:EC2_INSTANCE_LAUNCH
AccountId: 543038345491
AutoScalingGroupName: ASG-Sample-APP
AutoScalingGroupARN: arn:aws:autoscaling:us-east-1:543038345491:autoScalingGroup:cd936eeb-3543-4bae-8683-84dde6a70b89:autoScalingGroupName/ASG-Sample-APP
ActivityId: f6a636ae-97a9-71c3-7fd3-ccc662794ab6
Description: Launching a new EC2 instance: i-0b5aebb5f6303f0b2
Cause: At 2024-02-15T07:23:08Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 2. At 2024-02-15T07:23:11Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 2.
StartTime: 2024-02-15T07:23:13.884Z
EndTime: 2024-02-15T07:23:45.649Z
StatusCode: InProgress
StatusMessage:
Progress: 50
EC2InstanceId: i-0b5aebb5f6303f0b2
Details: {"Subnet ID":"subnet-03296ac2201e25af3","Availability Zone":"us-east-1b"}

Instance: i-0b5aebb5f6303f0b2 (APPASG)

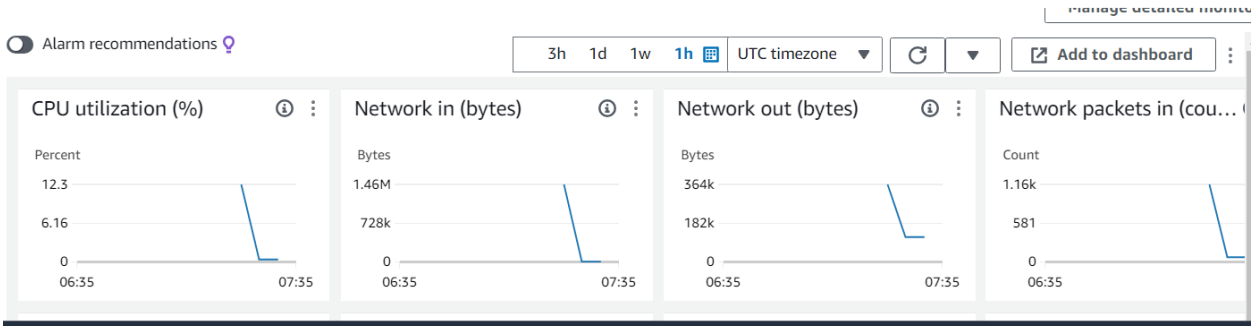

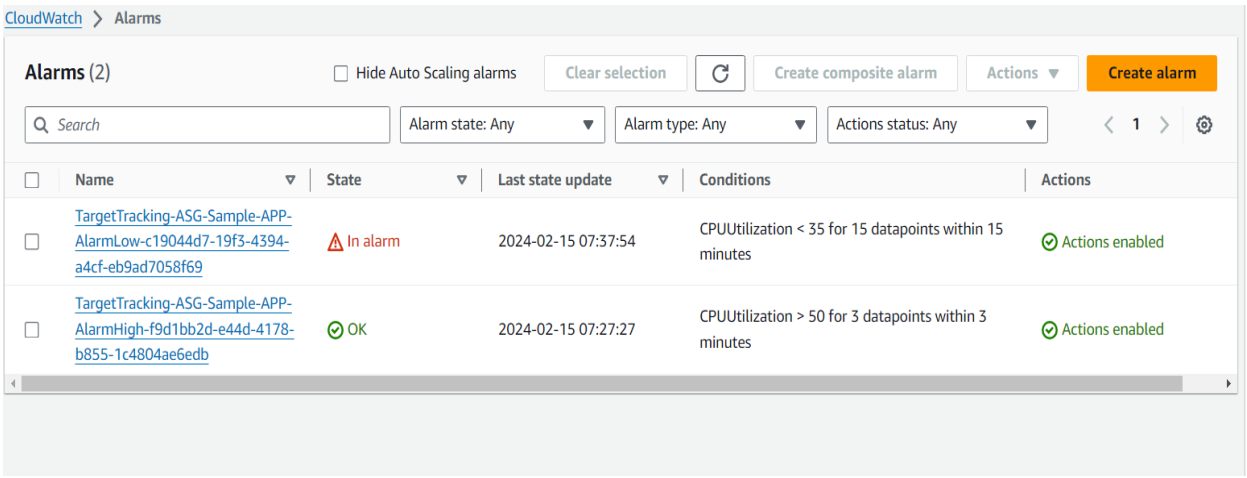
CloudWatch
Alarms
Alarms (2)
Hide Auto Scaling alarms
Clear selection
Create composite alarm
Actions
Create alarm
Search
Alarm state: Any
Alarm type: Any
Actions status: Any
1
Name
State
Last state update
Conditions
Actions
TargetTracking-ASG-Sample-APP-AlarmLow-c19044d7-19f3-4394-a4cf-eb9ad7058f69
In alarm
2024-02-15 07:37:54
CPUUtilization < 35 for 15 datapoints within 15 minutes
Actions enabled
TargetTracking-ASG-Sample-APP-AlarmHigh-f9d1bb2d-e44d-4178-b855-1c4804ae6edb
OK
2024-02-15 07:27:27
CPUUtilization > 50 for 3 datapoints within 3 minutes
Actions enabled